

Deep Adaptive Learning for Writer Identification based on Single Handwritten Word Images

Pattern Recognition Elsevier, vol.88, pp.64-74


November 2018

6th on Google Scholar Computer Vision and Pattern Recognition

IF 5.898

Main idea



- Explicit/Implicit information (use as much information as you can)



	Implicit: Writer identity: "bob"	Explicit: Word content: 'Imagine' Word length: 7 letters Character attribute: a, e, g, i, m, n
--	--	--

- Main/Auxiliary task (multi-task learning)
- Regularization?

CVL Dataset


- 310 Writers x 5 pages (English/German)
- Document, Line, Word level (two students manually)
- filter(min 20 instances, length>0) : 99,513 -> 28,735/70,788 (70,853/28,660)
- Image Preprocessing?



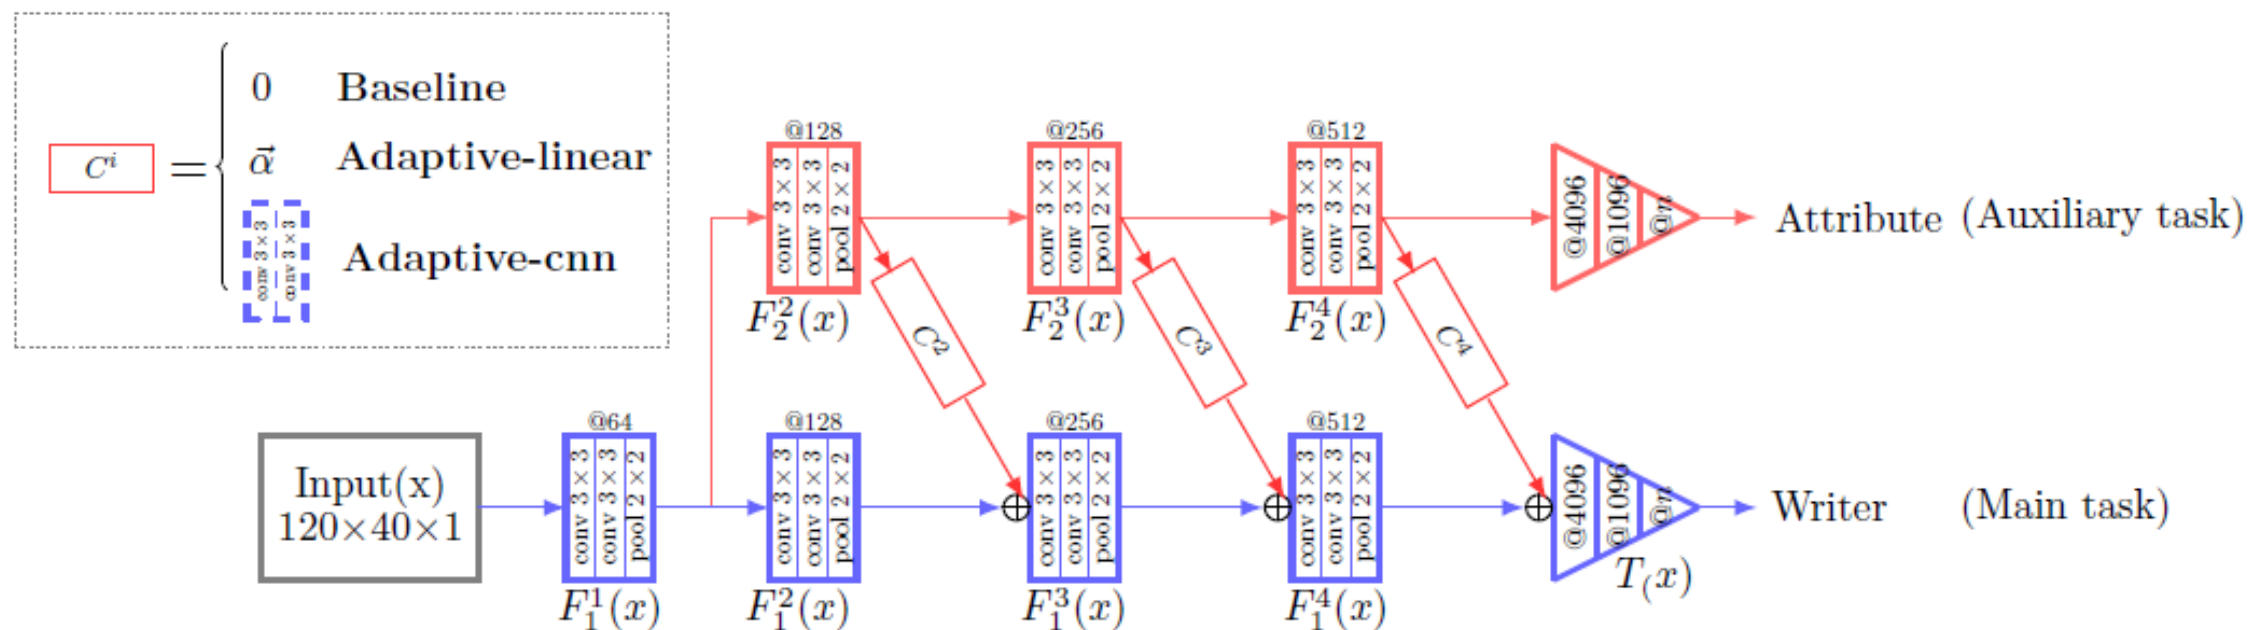
 $(4\theta, 12\theta)$ [illegible] $(4\theta, 12\theta)$ [illegible]

Proposed Method

- Single word images
- CNN for multi task learning (writer identification main task, auxiliary task?)
- Two pathways
- ~AlexNet, NIPS, 2012



Architecture



Auxiliary tasks

- Word recognition
- Character attribute recognition
- Word length estimation (1-13)

Model	Writer Identification			
	CVL		IAM	
	Top-1	Top-5	Top-1	Top-5
Baseline	75.3	92.5	66.0	82.9
Linear-adaptive	75.9	92.7	65.4	83.1
Deep-adaptive	79.1	94.3	68.3	85.2

	writer 1	writer 2	writer 3	writer 4	writer 5
word with 2 characters					
word with 3 characters					
word with 4 characters					
word with 5 characters					
word with 6 characters					
word with 7 characters					

Model?

- Main formula

$$in(F_1^{i+1}) = r(F_1^i) + C^i(r(F_2^i))$$

- Baseline


$$C^i(\cdot) = 0$$

- Linear

$$in_j(F_1^{i+1}) = \alpha_j \cdot r_j(F_1^i) + (1 - \alpha_j) \cdot r_j(F_2^i)$$

- Deep Adaptive

$$C^i(r(F_2^i)) = in(F_1^{i+1}) - r(F_1^i)$$



This is still seminary work

```
79  
80 O = Concatenate(axis=1, name='output_O')([T, R])
```

```
1 model2 = Model(inputs=x_input, outputs=[T, R, O], name='Model2')
```

```
1 def CustomLoss(y_true, y_pred):  
2     y_true1 = y_true[:, :MAX_WRITER]  
3     y_true2 = y_true[:, MAX_WRITER:]  
4  
5     y_pred1 = y_pred[:, :MAX_WRITER]  
6     y_pred2 = y_pred[:, MAX_WRITER:]  
7  
8     return alpha*K.categorical_crossentropy(y_true1, y_pred1)+(1-alpha)*K.categorical_crossentropy(y_true2, y_pred2)
```

```
10 alphaChanger = LambdaCallback(on_epoch_end=lambda epoch, _: K.set_value(alpha, 0.5+(epoch+1)*0.01))
```

```
1 model2.compile(optimizer=Adam(lr=0.0001), \  
2               loss={'output_T': categorical_crossentropy, 'output_R': categorical_crossentropy, 'output_O': CustomLoss}, \  
3               loss_weights={'output_T': 0, 'output_R': 0, 'output_O': 1}, \  
4               metrics=['acc'])
```

Implementation details

- Conv2D: kernel_size = (3, 3)
- MaxPool2D: pool_size = (2, 2)
- LeakyRelu(0.1)
- Dropout(0.5)
- cross-entropy loss (softmax activation/sigmoid activation?)
- Xavier initialization
- Adam(0.001)
- batch_size=100, n_iter=40 000 ~ 40 epoch! (My problem?)
- 83,589,227 params

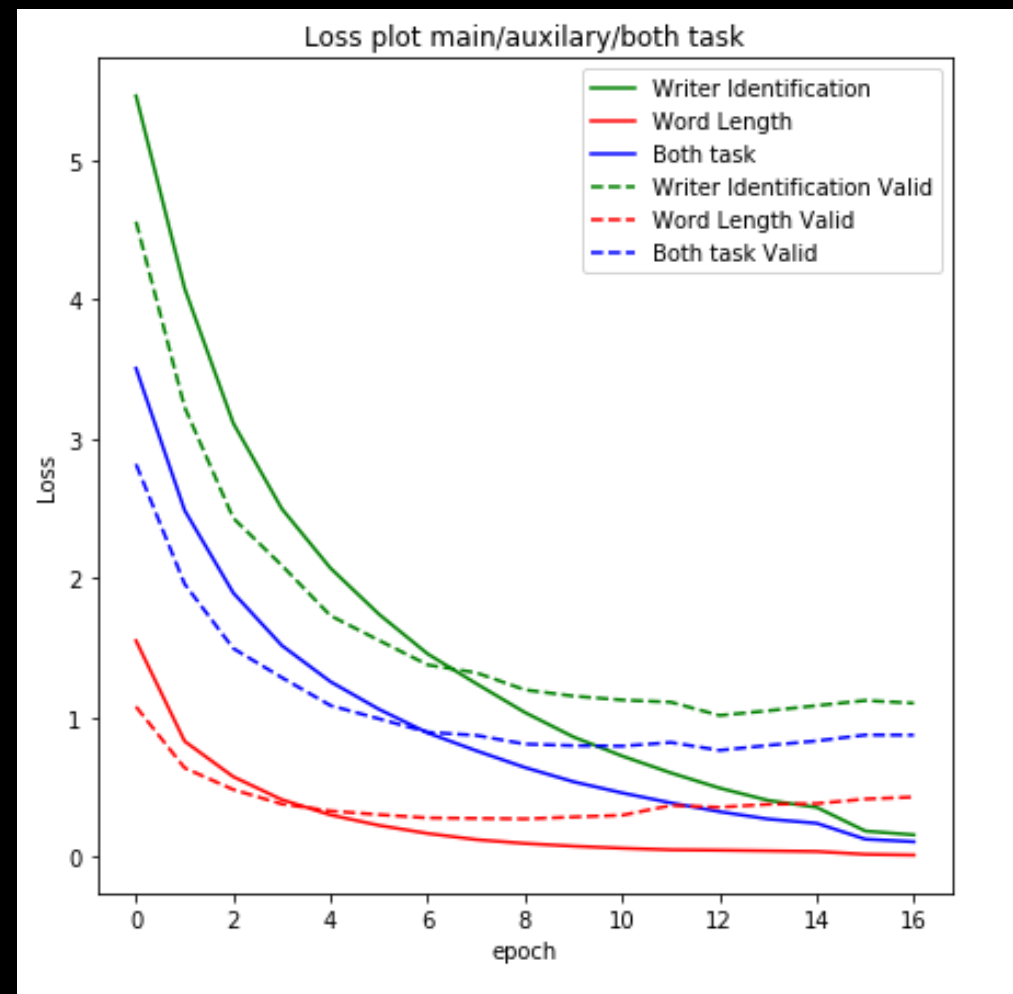
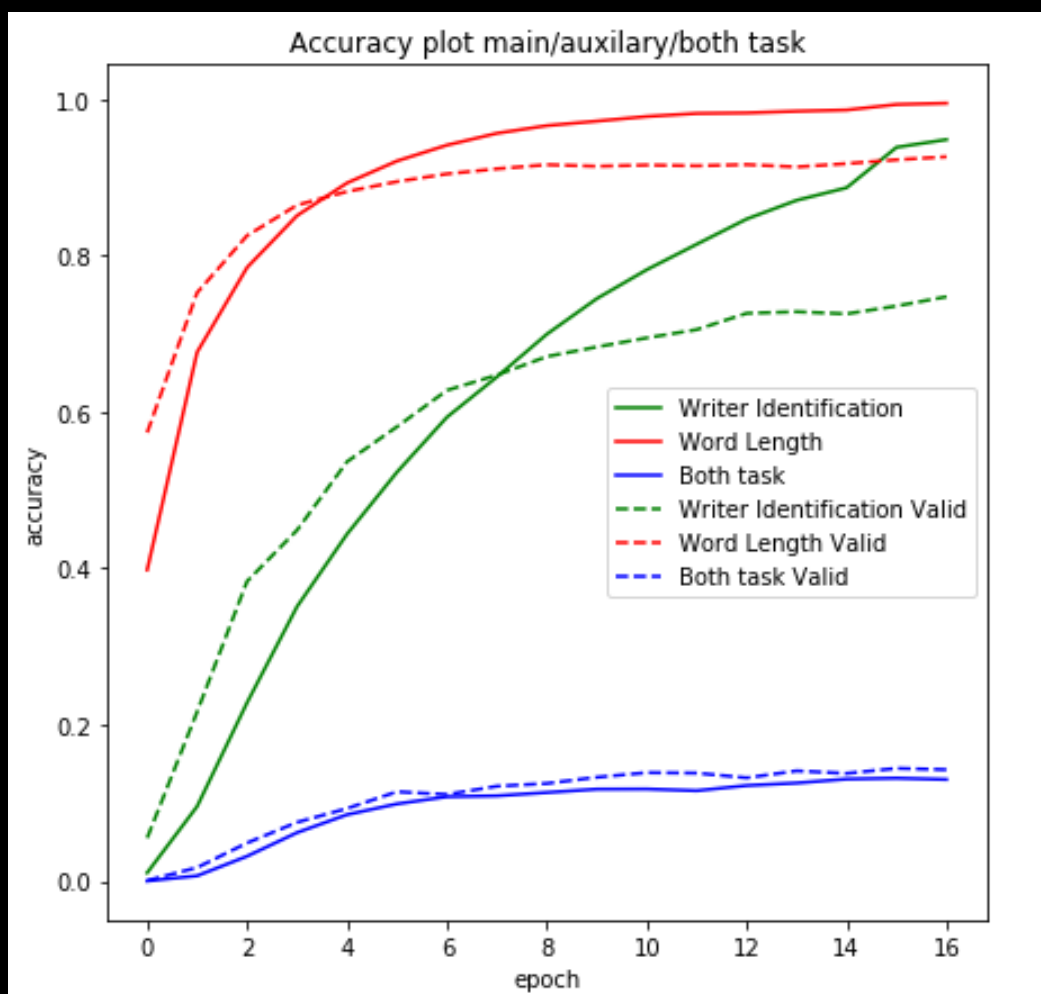
Training

- Combine two losses (O -> T, R)

$$\text{Loss}_{total} = (1 - \lambda)\text{Loss}_{au} + \lambda\text{Loss}_{wi}$$

- $\lambda = 0.5$ (10 epoch) + 0.066 (5 epoch) - increasing lambda effect?
- 0.896
- $\lambda = 0.5$ (0 epoch) + 0.01 (every epoch) - better one
- 0.9
- Weird lambda?

Early stopping? Overfitting? 15/40? Split?



Top1/Top5?

- Weird evaluation?
- Identification based on N word images from the same writer
- Randomly selected N word images for each writer
- Put them into CNN
- Use average response of the last CNN layer to recognize writer (x20 times, avg)

$$y = \frac{1}{N} \sum_i^N \text{CNN}(x_i)$$

Results (Paper)

Model	Input size: 40×120×1			
	CVL		IAM	
	W.I.	W.L.E.	W.I.	W.L.E.
Baseline	75.3	94.3	66.0	91.5
Linear-adaptive	75.9	92.7	65.4	90.4
Deep-adaptive	79.1	93.6	68.3	91.6
Training Time	8.5 hours			

Baseline: 68.62%

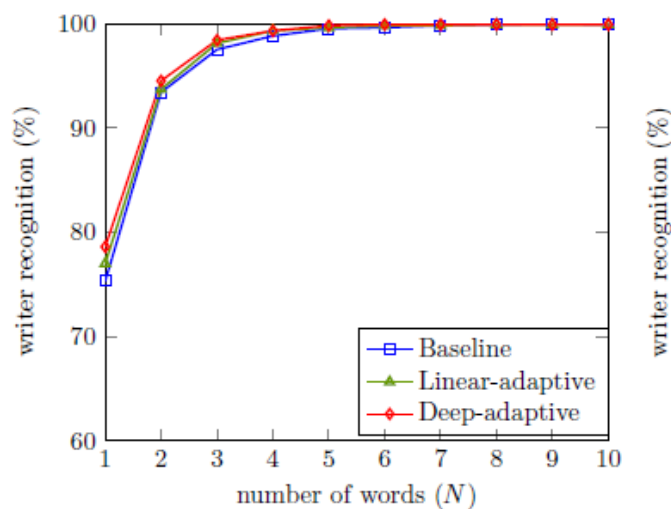
Deep-adaptive: 72.37%/92.02% (T/R)

Method	CVL	
	Top1	Top5
Hinge [1]	25.8	48.0
Quill [2]	29.4	52.6
Chain Code Pairs [3]	22.4	44.6
Chain Code Triplets [3]	28.8	51.4
COLD [34]	12.8	29.6
QuadHinge [28]	30.0	52.4
CoHinge [28]	25.9	46.9
CNN [6]	75.3	92.6
CNN+Adaptive	79.1	93.7

???

Table 2: Performance of writer identification using different adaptive learning methods with word recognition as the auxiliary task on the CVL and IAM datasets.

Model	Writer Identification				Word Recognition (<i>aux.</i>)			
	CVL		IAM		CVL		IAM	
	Top1	Top5	Top-1	Top-5	Top-1	Top-5	Top-1	Top-5
Baseline	75.3	92.4	65.7	83.5	95.1	99.1	93.5	98.7
Linear-adaptive	77.0	93.1	68.0	84.7	94.1	98.9	91.3	98.1
Deep-adaptive	78.6	93.7	69.5	86.1	94.5	99.0	92.6	98.4



(a) CVL data set

Figure 4: Performance (Top1) of writer identification using different numbers of words (from 1 to 10 words), using CNN models trained with word recognition as the auxiliary task on the CVL dataset (Figure (a)) and the IAM dataset (Figure (b)).

Misunderstanding?

From Fig. 4 we can see that writer-identification performance increases with more word images from the same writer. The Deep-adaptive model achieves the best results with different numbers of words for writer identification. The Top-1 performance for writer identification using the Deep-adaptive model was 79.1% and 68.3% when using one word, and this increases to 99.8% and 92.0% when using five words on CVL and IAM, respectively. For the special-

My improvements

- Next-To-Last layer
- Baseline:
 - Last: Top1/Top5/Top10: 67.90%/95.03%/98.87%
 - Next-To-Last: Top1/Top5/Top10: 66.99%/98.75%/99.73% (on Top1 is the small sample)
+3.72%
- Deep:
 - Last: Top1/Top5/Top10: 72.11%/97.20%/99.50% VS 79.10%/99.80%??
 - Next-To-Last: Top1/Top5/Top10: 70.53%/99.29%/99.98%
+2.09%

More ideas?

- Babysitting the learning process
 - (O? Better guidance through loss surface?)
- Preprocessing?
 - (Why n is the best auxiliary?)
- One word cropping?
 - (Augmentation some kind? Histogram with respect to n? ~Top3/5)
- Baseline (Hinton) to 75.3% and Next-To-Last = **Top5 Improvements publication!**

Questions?

Milan M. Čugurović

milan_cugurovic@matf.bg.ac.rs